Machine learning specialist task

<u>Description of the task:</u>
We want you to build a text classification service.
The model on which the service should be based on may be binary or multiclass/multilabel depending on the dataset of your choice.

<u>Expected artifacts:</u>
1. Model training code and description. It might be a script accompanied with description of data, training parameters and discussion on validation results OR Jupyter Notebook including everything aforementioned.
2. Web service code that has at least one endpoint '/predict' that accepts POST request with list of texts and returning list of corresponding classes.

<u>Additional info:</u>
The task description was intentionally made high-level so you are flexible in selecting dataset, frameworks, programming languages etc.
You can even use pretrained models but still please include data exploration and model validation analysis in the case.
It is highly recommended to use Docker for packaging the web service so we can run the application without dependencies hustle.
Expected completion time for the task is 4-8 hours (excluding model training time) depending on your experience with related technologies.